

PROBABILISTIC REASONING

Name : Guruprasath M R

Roll : CH.EN.U4AIE22015

LAB-2

Maximum Likelihood Estimation (MLE) for a normal distribution is given by:

$$\log L(\mu, \sigma^2) = -\frac{n}{2} \log(2\pi\sigma^2) - \frac{1}{2\sigma^2} \sum_{i=1}^n (x_i - \mu)^2$$

where μ is the mean, σ square is the variance, and x_i are the observed data points.

Explanation:

- The log-likelihood function for a normal distribution is defined. It calculates the log-likelihood based on the parameters `mu` (mean) and `sigma` (standard deviation). If `sigma` is non-positive, the function returns negative infinity to ensure valid parameter estimates.
- The sample mean and sample standard deviation are used as the initial guesses for the optimization process.
- The `minimize` function from `scipy.optimize` is used to find the parameter values that maximize the log-likelihood function. This is done by minimizing the negative of the log-likelihood function. The `bounds` argument ensures that `sigma` remains positive by setting a lower bound of `1e-6`.

Formula for MLE:

The Maximum Likelihood Estimates (MLE) for the mean μ and the variance σ square of a normal distribution are derived by maximizing the log-likelihood function.

For a sample of n observations x_1, x_2, \dots, x_n , the MLE for the mean and variance are given by:

$$\hat{\mu} = \frac{1}{n} \sum_{i=1}^n x_i$$

$$\hat{\sigma}^2 = \frac{1}{n} \sum_{i=1}^n (x_i - \hat{\mu})^2$$

```
In [ ]: import numpy as np
from scipy.optimize import minimize
import matplotlib.pyplot as plt

# Generate sample data from a normal distribution
np.random.seed(42)
sample_data = np.random.normal(loc=5.0, scale=2.0, size=100)
print("the sample data:\n", sample_data)
# Define the log-likelihood function for a normal distribution
def log_likelihood(params, data):
    mu, sigma = params[0], params[1]
    if sigma <= 0:
        return -np.inf # Return negative infinity if sigma is not positive
    n = len(data)
    return -((n / 2) * np.log(2 * np.pi * sigma ** 2) + (1 / (2 * sigma ** 2)) * np.sum((data - mu) ** 2))

# Initial guess for the parameters (mean and standard deviation)
initial_guess = [np.mean(sample_data), np.std(sample_data, ddof=1)]

# Perform the optimization to maximize the log-likelihood (minimize its negative)
result = minimize(lambda params: -log_likelihood(params, sample_data), initial_guess, \
                  method='L-BFGS-B', bounds=[(None, None), (1e-6, None)])

# The estimated parameters (mean and standard deviation)
mle_mean, mle_sigma = result.x

print(f"MLE for the mean: {mle_mean}")
print(f"MLE for the standard deviation: {mle_sigma}")

# Plot the data and the fitted distribution
plt.hist(sample_data, bins=20, density=True, alpha=0.6, color='g')

# Plot the fitted normal distribution
xmin, xmax = plt.xlim()
x = np.linspace(xmin, xmax, 100)
p = np.exp(-0.5 * ((x - mle_mean) / mle_sigma) ** 2) / (mle_sigma * np.sqrt(2 * np.pi))
plt.plot(x, p, 'k', linewidth=2)
title = "Fit results: mu = %.2f, sigma = %.2f" % (mle_mean, mle_sigma)
plt.title(title)
```

```
plt.show()
```

the sample data:

```
[ 5.99342831  4.7234714  6.29537708  8.04605971  4.53169325  4.53172609
 8.15842563  6.53486946  4.06105123  6.08512009  4.07316461  4.06854049
 5.48392454  1.17343951  1.55016433  3.87542494  2.97433776  5.62849467
 3.18395185  2.1753926  7.93129754  4.5484474  5.13505641  2.15050363
 3.91123455  5.22184518  2.69801285  5.75139604  3.79872262  4.4166125
 3.79658678  8.70455637  4.97300555  2.88457814  6.64508982  2.5583127
 5.41772719  1.08065975  2.3436279  5.39372247  6.47693316  5.34273656
 4.76870344  4.39779261  2.04295602  3.56031158  4.07872246  7.11424445
 5.68723658  1.47391969  5.64816794  4.22983544  3.646156  6.22335258
 7.06199904  6.86256024  3.32156495  4.38157525  5.66252686  6.95109025
 4.04165152  4.62868205  2.78733005  2.60758675  6.62505164  7.71248006
 4.85597976  7.0070658  5.72327205  3.70976049  5.72279121  8.07607313
 4.92834792  8.12928731 -0.23949021  6.64380501  5.17409414  4.4019853
 5.18352155  1.02486217  4.56065622  5.71422514  7.95578809  3.96345956
 3.38301279  3.99648591  6.83080424  5.65750222  3.94047959  6.02653487
 5.1941551  6.93728998  3.59589381  4.34467571  4.21578369  2.0729701
 5.59224055  5.52211054  5.01022691  4.53082573]
```

MLE for the mean: 4.792307004132914

MLE for the standard deviation: 1.8072324541860225

Fit results: $\mu = 4.79$, $\sigma = 1.81$

